

Project Report

Agentic Blog Writing System Using LangGraph

1 Introduction

This project implements a multi-agent autonomous blog generation system built using LangGraph.

Unlike traditional LLM applications that rely on a single prompt-response interaction, this system is architected as a state-driven multi-agent pipeline where multiple specialized agents collaborate to generate publication-ready blog content.

The system integrates:

- Strategic routing decisions
- Optional web-grounded research
- Structured task decomposition
- Parallel section generation
- AI-powered diagram generation
- Real-time UI streaming
- Downloadable Markdown outputs

The goal of this project is to demonstrate practical multi-agent orchestration using large language models and tool integrations.

2 Problem Statement

Standard LLM blog generators suffer from:

- No planning stage
- No research grounding
- Hallucinated citations
- Poor structural consistency
- No modular content generation
- No visual augmentation

This project solves these issues by introducing:

- Agent specialization
 - Structured outputs
 - Tool integration
 - Conditional execution
 - Parallelism
 - Deterministic merging
-

3 System Overview

The system follows this execution pipeline:

User Input (Topic, Date)



Router Agent



(Conditional)

Research Agent



Orchestrator Agent



Parallel Worker Agents



Reducer Subgraph

 |— Merge Content

 |— Image Planning Agent

 └— Image Generation Agent



Final Markdown + Images

The orchestration is handled by LangGraph, which manages state transitions between agents.

4 Core Architectural Components

4.1 LangGraph State

The system is built around a shared **State** object.

The state acts as the memory layer across all agents.

It includes:

- topic
- mode
- needs_research
- queries
- evidence
- plan
- sections
- merged_md
- image_specs
- final

Each agent:

- Reads from state
- Writes back to state
- Influences downstream execution

This design ensures coordination across multiple agents.

4.2 Router Agent (Strategic Decision Layer)

The Router Agent is the first node executed.

Its responsibilities:

- Analyze the topic
- Determine whether research is required
- Classify execution mode

The system defines three modes:

- ◆ **closed_book**

Used for evergreen topics that do not require recent information.

- ◆ **hybrid**

Used when topic is mostly conceptual but benefits from recent examples.

- ◆ **open_book**

Used for volatile topics such as:

- Latest news
- Weekly updates
- Policy changes
- Pricing updates

The router outputs:

- mode
- needs_research
- queries
- recency window

This decision controls the entire pipeline.

4.3 Research Agent (Tool-Integrated Agent)

Activated only when **needs_research = True**.

Responsibilities:

1. Call Tavily API
2. Fetch search results
3. Extract structured evidence using LLM
4. Deduplicate by URL
5. Apply recency filtering

In **open_book** mode, evidence is restricted to recent results (e.g., last 7 days).

This prevents outdated or hallucinated citations.

4.4 Orchestrator Agent (Planning Layer)

This agent performs task decomposition.

Instead of writing the blog directly, it generates a structured plan:

- Blog title
- Audience
- Tone
- Blog type
- 5–9 Tasks

Each Task includes:

- Section title
- Goal
- Bullet points
- Target word count
- Citation requirement
- Code requirement

This stage improves structure and coherence.

4.5 Parallel Worker Agents

For each task, a worker agent is spawned using LangGraph's `Send()`.

Each worker:

- Writes exactly one section
- Follows bullet instructions
- Meets word count constraints
- Cites only approved evidence URLs
- Includes code if required

Workers execute independently and in parallel.

This improves modularity and scalability.

4.6 Reducer Subgraph

After all workers finish, a reducer subgraph runs.

This subgraph consists of three internal nodes:

4.6.1 Merge Node

- Sorts sections by ID
 - Combines them into full Markdown
 - Adds blog title
-

4.6.2 Image Planning Agent

Reads full blog content and decides:

- Whether images are needed
- Maximum of 3 images
- Insert placeholders (e.g., [[IMAGE_1]])
- Create structured image specifications

This is an editorial reasoning step.

4.6.3 Image Generation Agent

For each image specification:

- Calls Gemini Image Model
- Saves image to `images/`
- Replaces placeholder with Markdown image tag

Graceful fallback is implemented in case image generation fails.

5 Frontend (Streamlit UI)

The frontend allows:

- Topic input
- As-of date selection
- Real-time execution streaming
- State summary display
- Evidence table
- Markdown preview with images

- Download Markdown
- Download image bundle
- Load past blogs

It interacts with the compiled LangGraph app.

6 Multi-Agent Characteristics

This project is multi-agentic because:

- It contains multiple autonomous decision-making units
- Each unit has a specialized responsibility
- Agents interact through shared state
- There is conditional branching
- Tool usage is integrated
- Parallel execution is implemented
- Execution is not linear

This is not prompt chaining.

It is orchestrated collaboration.

7 Technology Stack

- LangGraph – Multi-agent orchestration
 - OpenAI GPT-4.1-mini – Reasoning & writing
 - Tavily API – Web research
 - Gemini Image Model – Diagram generation
 - Streamlit – Frontend UI
 - Pydantic – Structured LLM outputs
-

8 Key Design Decisions

Structured Outputs

All major agents use Pydantic schemas to prevent malformed responses.

Conditional Execution

Research runs only when required.

Parallelism

Section writing is parallelized using LangGraph fan-out.

Recency Filtering

Open-book topics are restricted to recent evidence.

Tool Isolation

Workers only cite pre-approved evidence.

Graceful Failure Handling

Image generation failures do not break pipeline.

9 Strengths of the System

- Modular design
 - Scalable architecture
 - Strong hallucination control
 - Clean separation of concerns
 - Tool integration
 - UI transparency
 - Deterministic reducer logic
-

10 Limitations

- No reflection/critic loop
 - No persistent memory across sessions
 - No SEO optimization layer
 - No automated scheduling
 - No citation validation pass
-

11 Future Enhancements

- Add Critic Agent for quality scoring
- Add citation verification module
- Add autonomous publishing scheduler

- Add SEO optimization agent
 - Add long-term memory for writing style
 - Convert to API-based backend deployment
-

Conclusion

This project demonstrates a practical implementation of multi-agent orchestration using LangGraph.

It moves beyond simple prompt engineering and introduces:

- Decision agents
- Planning agents
- Tool-using agents
- Parallel workers
- Reducer subgraphs
- State-driven coordination

The system showcases how structured agent collaboration can produce high-quality, grounded, and modular AI-generated content.

It serves as a strong demonstration of:

- Agentic system design
 - LLM orchestration
 - Tool integration
 - Parallel execution
 - Production-style architecture
-